


Yifat Amir  
HW1  
Deep RL

Task	My Mean of Returns	My SD of Returns	Expert Mean of Returns	Expert SD of Returns
Ant	4852.50172309	74.2398030029	4782.16091724	132.555603119
Walker2d	4995.4266902	1242.65553931	5534.92655192	44.5002486121
Hopper	1722.43819117	1216.78097727	3777.9972139	3.20571253539
Reacher	-4.86192835166	1.69605609577	-3.89877452701	1.81825232762
Humanoid	10385.5954084	52.9418006081	10411.3670062	64.3865538357
HalfCheetah	4049.21748582	92.0304040342	4158.39699685	46.9716858278

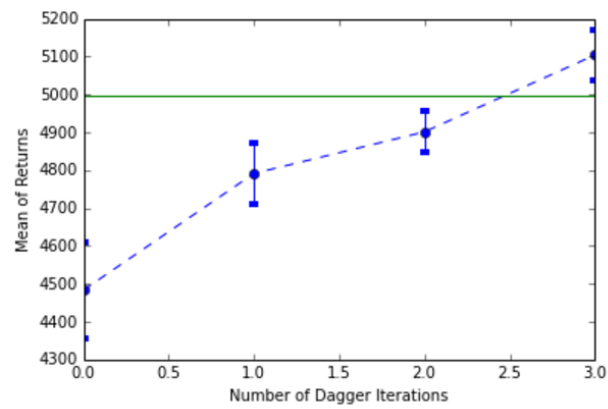
All tasks were run on a NN with two hidden layers, tanh activation, 20000 sample points, 20000 training iterations, batch size of 100, and learning rate of 0.001. Most tasks achieve comparable performance, except Hopper.



```
print creature
```

Ant

The graph above shows how performance varies with the value of the learning rate. I chose to experiment with this hyperparameter because it can greatly affect the optimization of the loss function. If the learning rate is too high, it is easy to overshoot a local optimum in the cost function. If the learning rate is too low, then it is hard to learn anything given a fixed number of training iterations.



```
print creature
```

Walker2d

The task above was run on a NN with two hidden layers, tanh activation, 20000 sample points, 10000 training iterations, batch size of 100, and learning rate of 0.001. There were 4 DAgger iterations.