

CS294 Deep Reinforcement Learning HW4: MPC Model-Based RL

Yifat Amir

```
In [1]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
```

Part 1

Fit a dynamics model to random data alone and use the learned dynamics model in your MPC controller to control the cheetah robot. Report your performance (copy/paste the log output into your report).

Iteration 0

AverageCost 121.468224013

StdCost 57.2858408949

MinimumCost 53.6009972585

MaximumCost 234.567366938

AverageReturn -66.0960621149

StdReturn 16.5012392242

MinimumReturn -98.4327335961

MaximumReturn -39.3160723071

Part 2

Run the full algorithm, including on-policy data aggregation, for 15 iterations. Make a graph of the performance (average return) at each iteration. How does performance change when the on-policy data is included?

I'm plotting only the first 6 iterations because my program is killed (9) consistently partway through iteration 7 (probably due to RAM). However, you can still see the improvement in the average return due to on-policy data.

```
In [2]: # python main.py -n 15
```

```
In [3]: returns = np.array([[-55.48267077, -80.45197874, -52.65621207, -79.43407
671, -62.87969339,
-98.4327336, -57.90643994, -39.31607231, -57.1266877, -77.27405592],
[ 367.6707783, 331.59849781, 340.56866406, 323.04083954, 335.528623
6,
363.92776312, 347.00587687, 264.78432737, 393.87546613, 345.505925
72],
[ 693.91343055, 896.65851939, 897.74801041, 910.33378461, 842.536152
05,
767.53420877, 901.6208748, 819.93068588, 822.1079186, 791.680131
74],
[ 973.90002293, 1004.31786705, 1037.80711735, 967.5827035, 933.3
3556119,
953.95272915, 1017.77948674, 895.31525909, 885.10568509, 879.9
8698304],
[ 824.88018926, 751.08901308, 820.39476689, 802.62657201, 812.547309
15,
857.3520731, 884.34524805, 874.87008172, 807.22266434, 862.945757
16],
[ 813.95781192, 816.13460029, 771.81553034, 883.36952685, 884.557410
11,
791.28056846, 800.65754863, 793.41804097, 792.37504429, 730.999163
84]])
```

```
In [4]: avg_returns = [np.mean(row) for row in returns]
print(avg_returns)
```

```
[-66.096062115000009, 341.35067625199997, 834.40637167999989, 954.90834
15130001, 829.82736747600006, 807.85652456999992]
```

```
In [5]: plt.plot(range(6), avg_returns)

plt.xlabel('Iteration')
plt.ylabel('Average Return')
plt.title('Half Cheetah MPC Returns')
plt.show()
```

